

# MACHINE LEARNING FOR DATA MINING

## LECTURE 2-2: NAIVE BAYES CLASSIFICATION

Siamak Sarmady (Urmia University of Technology)  
Sebastian Thrun, Katie Malone (Google)

### Naive Bayes method and Bayes Rule

- Bayes rule is the base of **probabilistic inference**.
- The rule was created by Rev. **Thomas Bayes**. He intended to use the principle to prove the **existence of God**.
- The word "**Naive**" refers to the **simplistic** nature of the algorithm.
- The rule has influenced **vast areas** of **science** including statistics and AI.

### Classification using Bayes' Rule

- In classification, Bayes' rule is used to **calculate** the **probabilities** of the **classes**.
- Based on the **training data** (seen items), we **select** a class that **minimizes** the expected **risk** (of selecting the wrong class)

## Probability

- **Coin toss:** If we had enough **knowledge** and **time**, we could predict the **exact outcome**.
  - Since we **do not have** enough knowledge and time, we use **probability** to make predictions.
- **Unobservable variables:** the piece(s) of knowledge that we **do not have access** to...
  - **In the case of coin:**
    - The only **observable** variable: the **outcome** of the toss (a dependent var too)
    - The **unobservable** variables: the **amount of force**, **angel** of force, the **point** of force
- Denoting the **unobservable** by  $z$  and the **observable** as  $x$ :

$$x = f(z)$$

where  $f(\cdot)$  is the **deterministic** function (such as mechanical equation) that defines the outcome **from** the **unobservable** pieces of knowledge.

## Minimizing Error

- Since we **don't have** enough knowledge (the **unobservable** variables) and we **cannot** model the process **this way**, we consider the outcome  $X$  as a **random** variable drawn from a probability distribution  $P(X=x)$ .
- The **outcome** of the coin toss is either head or tail denoted by  $X=1$  and  $X=0$  so:  $X \in \{1,0\}$
- These  $X$  are **Bernoulli distributed** where the parameter of the distribution  $p_0$  is the probability of the outcome **"head"**:

$$P_{(X=1)} = p_0 \quad \text{and} \quad P_{(X=0)} = 1 - P_{(X=1)} = 1 - p_0$$

- **Minimizing error:** if " $p_0 > 0.5$ " (e.g. 0.8), our **prediction** for the next toss will be **head**. That is because if we **choose** the **more probable** case, the **probability of error**, which is 1 minus the probability of our choice, will be **minimum**.
  - If the coin is **fair** ( $p_0=0.5$ ), we can choose **Heads every time** or use a fair coin to predict the outcome.

## Estimating from Samples

- If we **do not know** the probability **distribution**,  $P(X)$  and we want to **estimate** it from a given **sample**, then we are doing **statistics**.
- We have a **sample**  $X$ , from the probability distribution of the **observables** ( $x^t$  i.e. the outcome of toss  $t$ ). The **aim** is to build an **approximator** to it,  $\hat{p}(x)$ , using the sample  $X$ . The sample contains the outcome of previous  $N$  tosses.

$$\hat{p}_0 = \frac{\#\{\text{tosses with outcome "heads"}\}}{\#\{\text{tosses}\}}$$

- $x^t$  is 1 if the outcome of toss  $t$  is heads and 0 if it is tail
- Given the sample  $X=\{1, 1, 1, 0, 1, 0, 0, 1, 1\}$  the estimate will be

$$\hat{p}_0 = \frac{\sum_{t=1}^N x^t}{N} = \frac{6}{9}$$

## Classification (with Example)

- We want to **classify** customers as "**High risk**" and "**Low risk**" based on previous examples of people who **paid back** or **defaulted** a loan.
  - ▣ **Observable variables:** let's say we observe "**yearly income**" and "**savings**" of customers represented by two **random** variables  $X_1$  and  $X_2$  (and of course the outcome of loan payback)
  - ▣ **Non-observable variables:** other variables that could help us determine the class deterministically
- **The credibility of customer:** is a **Bernouli** random **variable**  $C$  (i.e. class) **conditioned** on **observable** variables  $X=[X_1, X_2]^T$ 
  - ▣  $C=1$  indicates a "High risk" customer
  - ▣  $C=0$  indicates a "Low risk" customer

## Classification

- If we know  $P(C|X_1, X_2)$  when a new application arrives with  $X_1=x_1$  and  $X_2=x_2$ , we can choose or predict the class using:

$$\begin{cases} C = 1, & \text{if } P(C = 1|x_1, x_2) > 0.5 \\ C = 0, & \text{otherwise} \end{cases}$$

- Or:

$$\begin{cases} C = 1, & \text{if } P(C = 1|x_1, x_2) > P(C = 0|x_1, x_2) \\ C = 0, & \text{otherwise} \end{cases}$$

- The probability of error is  $1 - \max(P(C = 1|x_1, x_2), P(C = 0|x_1, x_2))$
- If the observed variables are  $x=[x_1, x_2]^T$ , then the problem is to calculate  $P(C|x)$  using Bayes' rule:

$$P(C|x) = \frac{P(C) p(x|C)}{p(x)}$$

## Classification using Bayes' Rule

- **Prior probability:** is the knowledge (the probability estimate) we have of the value of  $C$  even before looking at the observable  $x$ .
  - $P(C=1)$  is called the prior probability that  $C$  takes the value 1, regardless of the  $x$  value.
  - The prior probabilities of two classes satisfy the following:

$$P(C=0) + P(C=1) = 1$$

- **Class likelihood:**  $p(x|C)$  is called the class likelihood and is the conditional probability that an event belonging to  $C$  has the associated observation value  $x$ 
  - In our case  $p(x_1, x_2|C=1)$  is the probability that a "high risk" customer has his  $X_1=x_1$  and  $X_2=x_2$ . This is obtained from the example data...
- **Evidence probability:**  $p(x)$  is the marginal probability that an observation  $x$  is seen, regardless of its class, i.e. positive or negative (in the example data)

$$p(x) = \sum p(x, C) = p(x|C = 1) P(C = 1) + p(x|C = 0) P(C = 0)$$

## Classification using Bayes' Rule

- **Posterior Probability:** combining the **prior** and **what data tells** us using Bayes' rule, we **calculate the posterior** probability of the concept,  $P(C|x)$ , **after seeing** the **observation x**

$$posterior = \frac{prior \times likelihood}{evidence}$$

- Since posterior is **normalized by** the **evidence**, the posteriors **sum up to 1**

$$P(C=0|x) + P(C=1|x) = 1$$

- After we have the **posteriors**, we can decide **which class** is more probable
- The **prior** and **likelihoods** are calculated by **counting** in the example data

## Multi-class Classification using Bayes' Rule

- In the general case, we have **K** mutually **exclusive classes**;  $C_i, i=1, \dots, K$
- For example in optical **digit** recognition the input is the bitmap image (pixel values), and there are **10 classes**. The probabilities satisfy:

$$P(C_i) \geq 0 \text{ and } \sum_{i=1}^K P(C_i) = 1$$

- $P(x|C_i)$  is the **probability** of **seeing x** as the input **when** it is known to belong to **class  $C_i$**
- The posterior probability of class  $C_i$  can be calculated as

$$P(C_i|x) = \frac{p(x|C_i)P(C_i)}{p(x)} = \frac{p(x|C_i)P(C_i)}{\sum_{k=1}^K p(x|C_k)P(C_k)}$$

## Multi-class Classification using Bayes' Rule

- For the **minimum error**, the Bayes' classifier chooses the class with the highest posterior probability; that is, we choose:

$$C_i \text{ if } P(C_i|x) = \max_k P(C_k|x)$$


i.e. the **k that maximizes**  $P(C_k|x)$

## Example 1:

- There is a **cancer type** who is observed in **1%** of the **population** i.e.  $P(c) = 0.01$
- There is a test for the cancer:
  - **If the person has cancer:** **90%** chance that test is positive (**sensitivity**).  
So there is **10%** chance that it won't be detected (**false negative**).
  - **If the person does not have cancer:** **90%** chance that the test is negative (**specificity**).  
So there is **10%** chance that it wrongly detects cancer (**false poitive**).
- **Question:** A person takes the test. The answer comes positive. How much is the probability that he has cancer?

### Example 1:


- ❑ If the person has cancer: 90% chance that test is positive (sensitivity).  
So there is 10% chance that it won't be detected.
- ❑ If the person does not have cancer: 90% chance that the test is negative (specificity).  
So there is 10% chance that it wrongly detects cancer.



### Example 1:


- ❑ **Question:** A person takes the test. The answer comes positive. How much is the probability that he has cancer?

$$0.9 / (9.9 + 0.9) = 8.33$$



## Bayes Rule

- Bayes rule:




- We have some prior probability. We then **incorporate some evidence** from a test, to arrive at posterior probability.
- In fact the above rule is calculated like this:

$$\text{Prior Probability} \cdot \text{Test Evidence} = \text{Posterior Probability}$$

## Bayes Rule

- Bayes rule:



Bayes rule (the above one is not normalized):

$$P(a|b) = \frac{P(a) \times P(b|a)}{p(b)}$$

Or

$$P(b|a) = \frac{P(b) \times P(a|b)}{p(a)}$$



## Example 1

- Prior:  $P(c) = 0.01$  (1%)  
 $P(\text{Pos} | c) = 0.9$  (90%)  
 $P(c') = 0.99$   
 $P(\text{Neg} | c') = 0.9 \Rightarrow P(\text{Pos} | c') = 0.1$

Bayes rule:  
 $P(a | b) = P(a) \cdot P(b | a)$

Normalizer:  $P(b)$

$$1. P(c | \text{Pos.}) := P(c) \cdot P(\text{Pos.} | c) = 0.01 \cdot 0.9 = 0.009$$

but this is not quite right. It is not normalized. In order to normalize we should also calculate:

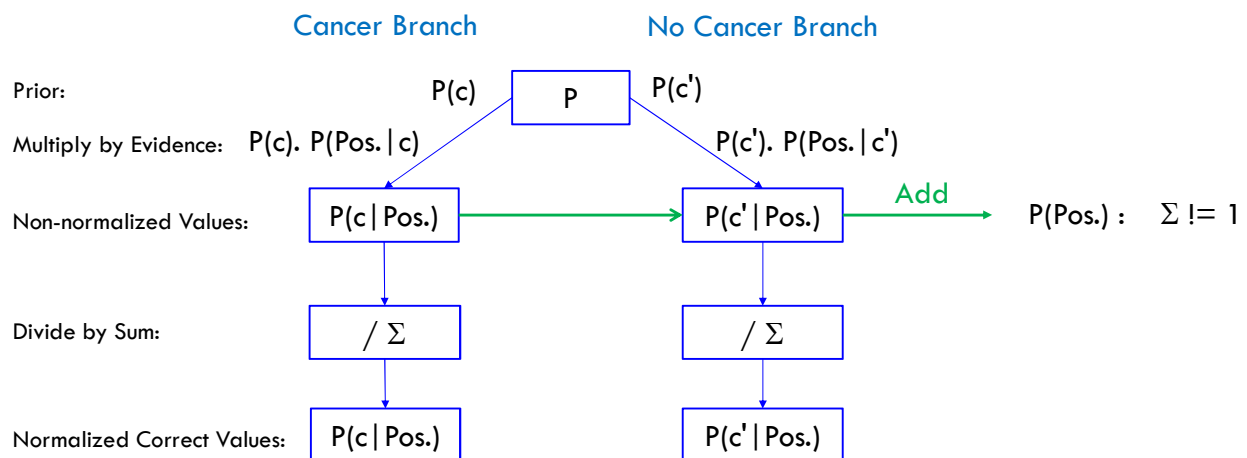
$$2. P(c' | \text{Pos.}) := P(c') \cdot P(\text{Pos.} | c') = 0.99 \cdot 0.1 = 0.099$$

The problem is that the above 1 & 2 do not add up to 1. So we need to normalize them to 1.

- **First:** We add them, the sum is 0.108, and this is in fact  $P(\text{pos}) = P(c | \text{Pos.}) + P(c' | \text{Pos.})$  or the normalizer.
- **Second:** We find correct values for the above probabilities (which add up to 1):
  - $P(c | \text{Pos.}) = 0.009 / 0.108 = 0.0833$
  - $P(c' | \text{Pos.}) = 0.099 / 0.108 = 0.9167$


## Example 1 (algorithm for bayes rule):

- $P(c) = 0.01$  Prior
- $P(\text{Pos.} | c) = 0.9$  Sensitivity
- $P(\text{Neg.} | c') = 0.9$  Specificity



## Example 2: Text Learning (Learning from Documents) using Naive Bayes


- Suppose we have **large number of emails** from two person, Chris and Sarah. For the simplicity assume these emails contain **only 3 type** of words (in reality it could be 30,000).
- The **difference** between these two series of emails is the **frequencies** of the words. They use these words with different frequencies.
- Naive Bayes allows to **guess** who is **more likely to be** the writer of an email.



- Chris likes to talk **more about** "Deal" while Sarah **mostly likes** to talk about "Love"...

### Example 2-1:

- Assume these priorities:  $P(\text{Chris}) = 0.5$ ,  $P(\text{Sarah}) = 0.5$





- We have received an email containing this text: "Love Life"
- Who is more likely to be the sender?
- This is relatively easy to guess but...

### Example 2-2:

Bayes rule:  
 $P(a|b,c) = P(a) \cdot P(b|a) \cdot P(c|a)$   
 Normalizer:  $P(b,c)$

- Priors:  $P(\text{Chris}) = 0.5$ ,  $P(\text{Sarah}) = 0.5$



- Now assume we have received another email: "Life Deal"
  - Who is more likely to be the sender this time? Not as easy as previous case.
  - $P(\text{Chris} | \text{Life}, \text{Deal}) := P(\text{Chris}) \cdot P(\text{Life} | \text{Chris}) \cdot P(\text{Deal} | \text{Chris}) = 0.5 \cdot 0.1 \cdot 0.8 = 0.04$
  - $P(\text{Sarah} | \text{Life}, \text{Deal}) := P(\text{Sarah}) \cdot P(\text{Life} | \text{Sarah}) \cdot P(\text{Deal} | \text{Sarah}) = 0.5 \cdot 0.2 \cdot 0.3 = 0.03$
- (so higher probability that Chris sent the letter)

### Example 2-2:

- However notice that above probabilities are not normalized (do not sum up to 1) and therefore cannot be treated like probability. In order to get correct probabilities we need to normalize them.
- $P(\text{Chris} | \text{Life}, \text{Deal}) = 0.04$
- $P(\text{Sarah} | \text{Life}, \text{Deal}) = 0.03$
- $\Sigma = 0.04 + 0.03 = 0.07$
- $P(\text{Chris} | \text{Life}, \text{Deal}) = 0.04 / 0.07 = 0.571$
- $P(\text{Sarah} | \text{Life}, \text{Deal}) = 0.03 / 0.07 = 0.428$
- Above values are proper probabilities now.

### Example 2-3:

- Now, let's assume we received another email with the text "Love Deal". Who is the likely sender.
- Non-normalized Probabilities:
  - $P(\text{Chris} | \text{Love, Deal}) = 0.5 * 0.1 * 0.8 = 0.04$
  - $P(\text{Sarah} | \text{Love, Deal}) = 0.5 * 0.5 * 0.2 = 0.05$
- $\Sigma = 0.04 + 0.05 = 0.09$
- Normalized correct probabilities:
  - $P(\text{Chris} | \text{Life, Deal}) = 0.04 / 0.09 = 0.444$
  - $P(\text{Sarah} | \text{Life, Deal}) = 0.05 / 0.09 = 0.555$
- So the new email is likely to be sent by Sarah

### Naive Bayes

- So, to sum up




- **A:**  $P_{(A)} \cdot P_{A2} \cdot P_{A3}$  this is in fact non-normalized of  $P(A | F2, F3)$
- **B:**  $P_{(B)} \cdot P_{B3} \cdot P_{B5}$  this is in fact non-normalized of  $P(B | F3, F5)$
- What comes out of above calculations, is the ratio between the probabilities of Label A and Label B
- If we need probabilities, we need to normalize them (i.e. Divide each to the sum of both).
- So Naive Bayes lets you identify the label (class) which is more likely, by giving it the features (evidence)

## Naive Bayes – Why Naive

- Naive Bayes is a **powerful** method. You can identify the **writer of a book, email or text** using it.
- One of the reasons it is **called** "Naive" (simplistic) is that, it **ignores** one of below things. Guess which one:
  - Words
  - Word Orders
  - Length of Message
- The probability calculation **ignores the order of words**. If you change the order of words, the meanings change or the text may not make sense at all.
- So the method **considers occurrence and frequencies** of words, **but not the order** of them. But still it works well in some situations.
- **Another reason** is that it always selects by minimizing risk.

## Naive Bayes

- Advantages:
  - It is **easy to implement** and **efficient**.
  - Can work with a really **big feature space** (e.g. 20000 words)
- Disadvantages:
  - It can break in funny ways. For example when google initially started to work, when people **searched for** "Chicago Bulls" which is a basketball team, google **would show** results about "Bulls" (cows) and the city of "Chicago". So things sensitive to order of the features are not very suitable for Naive Bayes method.



=



or

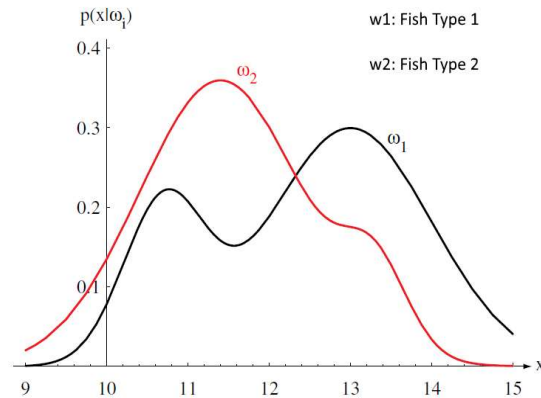


### Example 3 – Separating fishes based on their length

- An automatic machine should separate 2 types of fish on a conveyer belt. Probability distribution of each fish being of different lengths has been provided in below graph.
- If we know that 70% of fishes are of type 1 and 30% of type 2 (priori probability)
  - ▣ if the length of a fish is 11.5, what is the type?
  - ▣ If the length is 13, what is the type?

#### Assumptions:

- ▣  $P(11.5 | w_1) = 0.17$
- ▣  $P(11.5 | w_2) = 0.35$
- ▣  $P(13 | w_1) = 0.3$
- ▣  $P(13 | w_2) = 0.2$



### Example 3 – Separating fishes based on their length

- $P(w_1 | 11.5) = P(w_1) * P(11.5 | w_1) = 0.7 * 0.17 = 0.12$
- $P(w_2 | 11.5) = P(w_2) * P(11.5 | w_2) = 0.3 * 0.35 = 0.105$
- So the probability of a 11.5 cm fish of being a Type 1 is higher.
- $P(w_1 | 13) = P(w_1) * P(13 | w_1) = 0.7 * 0.3 = 0.21$
- $P(w_2 | 13) = P(w_2) * P(13 | w_2) = 0.3 * 0.2 = 0.06$
- Again the probability of a 13 cm fish of being a Type 1 is higher.

### Example 4 – Spam Email Identification

- Assume 10% of all emails are spam (priori probability).
- Also assume the distribution of words in spam and non-spam emails are as follows (the actual distribution has much more words)
  - ▣ Spam: Rolex 0.4, Winner 0.4, Send 0.2
  - ▣ Non-spam: Send 0.4, Exercise 0.3, Winner 0.2, Rolex 0.1
- If we have received 2 emails with the following word distributions, which one is likely to be spam:
  - ▣ Email 1: Rolex winner
  - ▣ Email 2: Send winner

### Example 4 – Spam Email Identification

- Email1:
  - ▣  $P(\text{spam} | \text{Rolex, winner}) = P(\text{spam}) * P(\text{rolex} | \text{spam}) * P(\text{winner} | \text{spam})$   
 $= 0.1 * 0.4 * 0.4 = 0.016 \rightarrow \text{Normlaized} = 0.47$
  - ▣  $P(!\text{spam} | \text{Rolex, winner}) = P(!\text{spam}) * P(\text{rolex} | !\text{spam}) * P(\text{winner} | !\text{spam})$   
 $= 0.9 * 0.1 * 0.2 = 0.018 \rightarrow \text{Normlaized} = 0.52$
- Email2:
  - ▣  $P(\text{spam} | \text{send, winner}) = P(\text{spam}) * P(\text{send} | \text{spam}) * P(\text{winner} | \text{spam})$   
 $= 0.1 * 0.2 * 0.4 = 0.008 \rightarrow \text{Normlaized} = 0.1$
  - ▣  $P(!\text{spam} | \text{send, winner}) = P(!\text{spam}) * P(\text{send} | !\text{spam}) * P(\text{winner} | !\text{spam})$   
 $= 0.9 * 0.4 * 0.2 = 0.072 \rightarrow \text{Normlaized} = 0.9$

## Training and Testing Classifiers

### Training and Test Data Sets

- Over fitting

### Training and Testing Classifiers

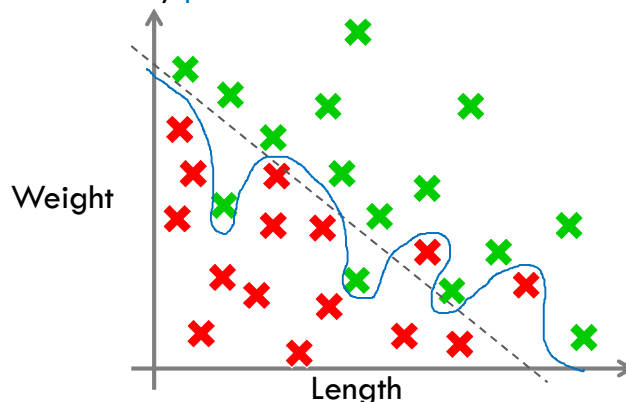
- We do **training** and **evaluation** (testing) on two separate data sets. Using an **independent "test set"** allows to test and evaluate the performance **independent** from **training data**.
- 1. Randomly **divide the data** (features and labels) to two groups, namely **"training set"** and **"test set"**.
  - The training set will normally be **80%-90%** of data items while test set could be **20%-10%** of them.
- 2. Train the classifier with "training set".
- 3. Use the features of the "test set" to predict labels using classifier.
- 4. Now compare the predicted values and the test set labels.
- 5. Calculate the **percentage** of **correct predictions**. The percentage of the correct predictions, determines the performance of the classifier.



## Overfitting

**Over-fitting:** sometimes we might over-fit the classifier to the training data. In that case the classifier will **no more generalize** properly.

- The classifier has in fact **memorized** the training set instead of finding the patterns in it.
- It will get 100% accuracy with the training data but it has no clue **how to generalize** to new data. It will have very **poor results** with other **unseen data** (like test data set).



## Conclusion

- As we saw, we **should not look** at classifiers as **black boxes**.
  - We should **understand how** they work, so that we can decide whether they are suitable for a problem.
  - In addition knowledge of **inner works** of a method helps in **adjustment of settings** and **better use** of the method and therefore **better performance**.
  - And also we should **test and see** how a **specific classifier** works for a specific problem.
  - Testing **should be done** with a **separate** data set which was not seen by the classifier during test phase.

## Python and Scikit-Learn Library

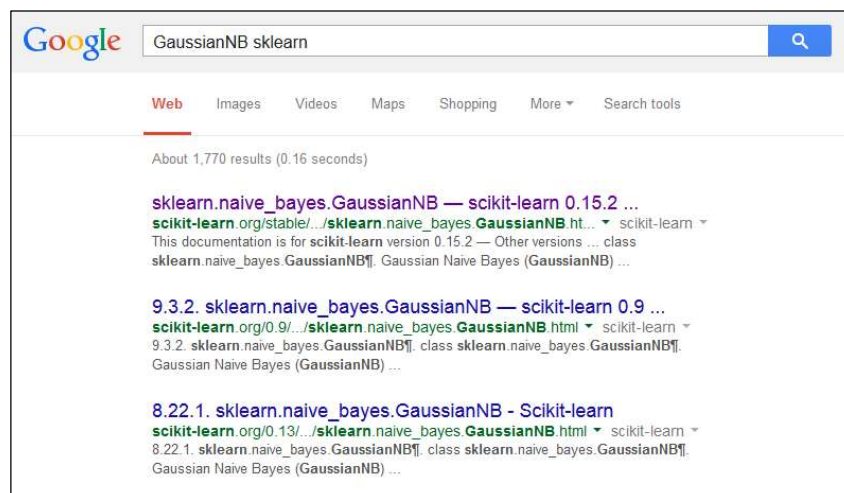
### Using Python and Scikit-Learn library

- We use Python version 3.5 because most of the libraries in the field of machine learning have become available for this version.
- We should either install Python and then add sickit-learn (and a few other libraries) to it, or use a Python-Bundle which puts together python and several libraries.
- We will use Winpython which includes several useful libraries...
- The library which we will use for Machine-Learning is called Scikit-Learn

## Using Python and Scikit-Learn library

- The website of Scikit-Learn is: <http://scikit-learn.org>
- The documentation can be found at: <http://scikit-learn.org/stable/documentation.html>
- Scikit-Learn is also known as "sklearn".
- If you want examples of a specific function just search for "sklearn" and the name of the function.

## Using Python and Scikit-Learn library



## Using Python and Scikit-Learn library

The screenshot shows the Scikit-Learn documentation page for the `GaussianNB` class. The page layout includes a top navigation bar with links for Home, Installation, Documentation, and Examples. A search bar is located on the right. The main content area is titled `sklearn.naive_bayes.GaussianNB`. It lists the class as `sklearn.naive_bayes.GaussianNB` and provides a detailed description of its attributes: `class_prior_` (array, shape = [n\_classes]), `theta_` (array, shape = [n\_classes, n\_features]), and `sigma_` (array, shape = [n\_classes, n\_features]). The page also includes an 'Examples' section with a code snippet demonstrating how to use the `GaussianNB` class for classification. The code snippet is as follows:

```
>>> import numpy as np
>>> X = np.array([[-1, -1], [-2, -1], [-3, -2], [1, 1], [2, 1], [3, 2]])
>>> Y = np.array([1, 1, 1, 2, 2, 2])
>>> from sklearn.naive_bayes import GaussianNB
>>> clf = GaussianNB()
>>> clf.fit(X, Y)
>>> print(clf.predict([[-0.8, -1]]))
[1]
```

Below the examples, the 'Methods' section lists the following methods:

- `fit(X, y)`: Fit Gaussian Naive Bayes according to X, y
- `get_params(deep=True)`: Get parameters for this estimator
- `predict(X)`: Perform classification on an array of test vectors X

## Sklearn - Example 1

- The example code is as follows:

```
import numpy as np

features_train = np.array([[-1, -1], [-2, -1], [-3, -2], [1, 1], [2, 1], [3, 3]])
labels_train = np.array([1, 1, 1, 2, 2, 2])
features_test = np.array([[-2, -2], [-2, -3], [2, 3], [1, 2]])
labels_test = np.array([1, 1, 1, 2])

from sklearn.naive_bayes import GaussianNB
clf = GaussianNB()
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)

print("Test labels: ", labels_test)
print("Predicted labels: ", pred)

from sklearn.metrics import accuracy_score
print("Accuracy: ", accuracy_score(pred, labels_test))

print("\nPredicted label for ", [-0.8, -1], " is ", (clf.predict([[-0.8, -1]])))
```

## Sklearn - Example 1


- Put the example code in a file with ".py" extension and run the code this way:

```
F:\mc-code\nb1>python nb1 1.py
Test labels:      [1 1 1 2]
Predicted labels: [1 1 2 2]
Accuracy:         0.75

Predicted label for [-0.8, -1] is [1]
F:\mc-code\nb1>
```


## Sklearn - Example 2

- What we are going to do in the 2nd example, is to take a training data set which contains a few hundred data points for the car speed adjustment example.




## Using Scikit Learn Python Library

- We then try to find a decision boundry for the data (i.e. Learning phase)



## Using Scikit Learn Python Library

- The classes painted with different colors



## Using Scikit Learn Python Library

- We can then use the boundary to decide which class an unseen arbitrary data belongs to (generalization step).

